

Part - b

6

- (b) → When data in tables are stored across different sites, there is a high chance of data redundancy to occur. Data redundancy occurs when the same piece of data exists in multiple places, whereas data inconsistency is when same data exists in different formats in multiple tables. Data redundancy can cause data inconsistency, which can provide a company with unreliable and/or meaningless information.
- If data consistency is not followed, then there is a high chance of data corruption due to result of overwriting & overreading the same data again & again. If file gets corrupt, then the queries to retrieve the data won't work properly.
- Data inconsistency also increases the database size & the complexity of the database. This is due to the repetition of the same data stored across multiple websites again & again.
- Data inconsistency also will end up increasing the cost, as more cost would be required for a bigger database. To maintain such a large database is also a tough task.

To avoid these problems & maintain data consistency, we can follow the following methods :-

### (i) Master data :-

If it is a single source of common business data that is shared across several applications or systems. Although master data does not reduce the occurrences of data redundancy, it allows the company to work around & accept a certain level of data redundancy. This is because the use of master data ensures that in the event a data piece changes, an organisation only needs to update one piece of data. In this case, redundant data is consistently updated & provides the same info.

### (ii) Database normalization :-

- Database normalization is the process of efficiently organising data in a database so that redundant data is eliminated. This process can ensure that all of a company's data looks & reads similarly across all records. By implementing

data normalization, an organization standardizes data fields such as customer names, addresses, & phone numbers.

- Normalizing data involves organising the columns & tables of a database to make sure their dependencies are enforced correctly. The "normal form" refers to a set of rules on normalizing data, & a database is known as 'normalized' if it's free of delete, update & insert anomalies.

Requirements for managing data consistency:-

- High speed, low <sup>input</sup> data comparison.
- Support for heterogeneous databases.
- Capability for handling large data volumes.
- Data security.
- Easy to use, understand, configure, deploy & diagnose.
- Low impact on hardware & network resources.

Challenges in maintaining data consistency:-

→ Migration errors:-

Different kinds of migration tools are employed to facilitate the initial load of the target database before replication can begin. Differences in config for handling data by the migration tools & replication products can result in data discrepancies.

→ Left & shift workload to cloud:-

Since the world is moving towards cloud, the left & shift of database workload from on-premises to cloud is need of today's IT world.

→ Differences in source & target:-

Differences in source & target database config, for example diff. encoding, locales, endianness or database versions can cause subtle discrepancies to happen during migration & replication.

→ Instantaneous errors:-

Before migration or replication can begin, the target database will need to be instantiated with correct schema & constraints. Failure to do so will result in source & target being out of sync.

Part - C

- 8) b) → Distributed lock manager can be used for this scenario. In this approach, functionality of locking is implemented by lock managers at each site.

- Lock manages control access to local data items.
- Locking is performed separately on each site accessed by transaction.
  - \* Every replica must be locked & updated
  - \* But special protocols may be used for replicas.

### Deadlock handling

Consider the following 2 transactions in history, with item X in transaction  $T_1$  at site 1, & item Y in transaction  $T_2$  at site 2:

$T_1$	wait(X) write(Y)	$T_2$	wait(X) write(Y)
X-lock on X write(X)	--		
		X-lock on Y write(Y) wait for X-lock on X	

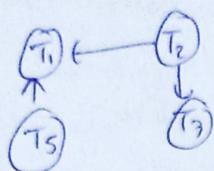
Wait for X-lock  
on Y

Result: Deadlock which can't be detected locally at either site.

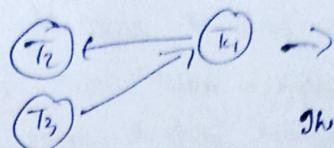
### Deadlock detection

- In the centralised deadlock-detection approach, global wait-for-graph is constructed & maintained in a single site:
  - Real graph: Real, but unknown state of system.
  - Constructed graph: Approximation generated by controller during execution of its algorithm.
- The global wait-for graph can be constructed when:
  - a new edge is inserted in or removed from one of local wait-for graphs.
  - a no. of changes have occurred in a local wait-for graph.
  - the coordinator needs to invoke cycle-detector.
- If coordinator finds a cycle, it selects a victim & notifies all sites. The sites roll back the victim transaction.

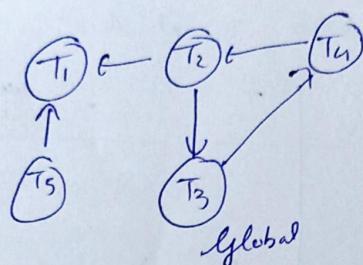
Local & global wait - for graphs



Sit S<sub>1</sub>



Sit S<sub>2</sub>

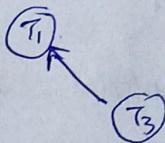


wait - for graph for false cycles

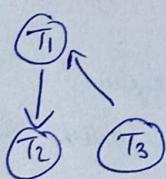
Initial state:



S<sub>1</sub>



S<sub>2</sub>



coordinator

False cycles:-

- Suppose that starting from the state shown above.
  - T<sub>2</sub> releases resource at S<sub>1</sub>.  
→ resulting in a message from Transaction manager at site S<sub>1</sub> to coordinator.
  - T<sub>2</sub> requests a resource held by T<sub>3</sub> at site S<sub>2</sub>  
→ resulting in a message from Transaction manager at site S<sub>2</sub> to coordinator.
- Suppose further that insert message reaches before delete message.  
→ This can happen due to network delays.

This fig. depicts a system of 2 sites, each maintaining its local wait-for graph. Transactions T<sub>1</sub>, T<sub>2</sub> appear in both graphs, indicating that transactions have requested items at both sites. When a transaction T<sub>i</sub> on site S<sub>1</sub> needs a resource in site S<sub>2</sub>, it sends a req msg to site S<sub>2</sub>. If resource is held by transaction T<sub>j</sub>, system inserts an edge T<sub>i</sub> → T<sub>j</sub> in local wait-for graph of site S<sub>2</sub>.

global

In centralised deadlock detection approach, the system constructs & maintains a global wait-for graph in a single site: the deadlock-detection coordinator. Since there is communication delay in the system, we must distinguish b/w 2 types of wait-for graphs.

- The coordinator would soon find a false cycle  
 $T_1 \rightarrow T_2 \rightarrow T_3 \rightarrow T_1$
- The false cycle above never existed in reality.
- False cycles cannot occur if 2-phase locking is used.

### Distributed deadlocks:

- Unnecessary rollbacks may result
  - when deadlock has indeed occurred & a victim has been picked, if meanwhile one of transactions was aborted for reasons unrelated to no deadlock.
  - Due to false cycles in global wait-for graph, however, likelihood of false cycle is low
- In distributed deadlock-detection approach, sites exchange wait-for info & check for deadlocks
  - Expensive & not used in practice.

Q) → Torrenting is, simply put, a file-sharing technology. Users, peers in this case, connect & share the files online. When you are using torrents you don't download files from a single website / source. Files circulate around the peers who seed them. It means that you download stuff from one direct source but you also download chunks from other users of same torrent. That enables a smooth transfer that's hard to trace.

- Peers are users who are involved in the sharing process through torrent P2P network. You are a peer as long as you keep sharing file on the given network.
- Seeders are users who are downloading but also uploading the file for other users to download. When you download something through torrent, you download certain amounts of data but you also upload a bit. That upload means that other users will get chunk they need when you upload it.
- Indexes are websites that work as search engines for files & content that can be downloaded through torrents.

Scalability is used to enable the working of torrent flow by large numbers of torrent due to access from multiple geographical locations.

- Torrent protocol enables decentralization of its resources by making use of

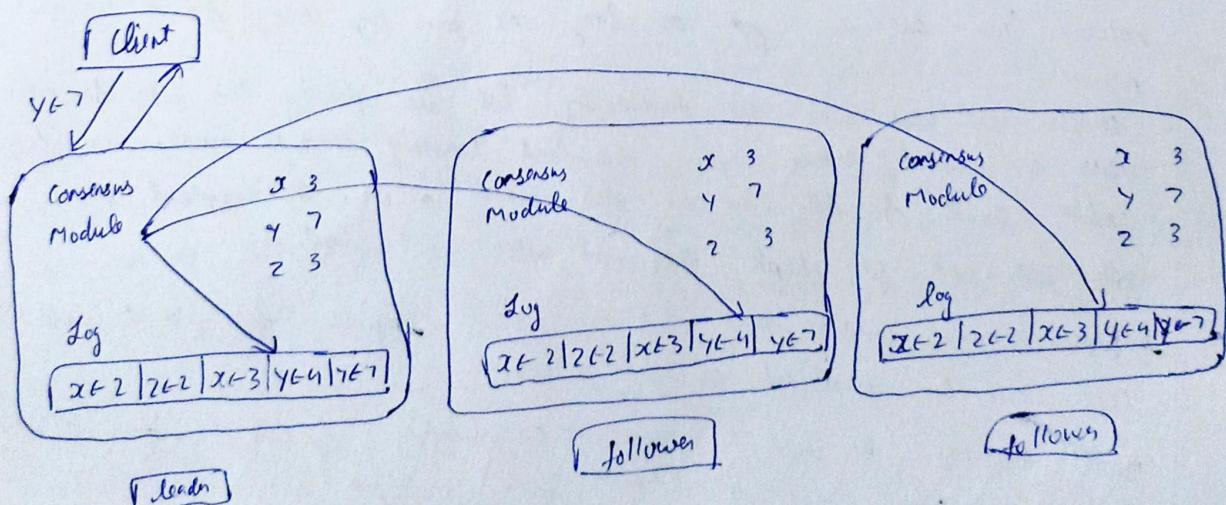
peer-to-peer network.

- A small torrent file is created to represent a file or a folder to be shared.
- Users can download the required files using a unique magnet link associated to each file or torrent.
- To learn the internet location of peers, we use distributed hash tables. The technology to ensure fault tolerance is fault-tolerant services using replicated state machines.

### Fault-Tolerant Services using replicated State machines

- Key requirement: makes the torrent service fault tolerant. Eg: lock manager, key-value storage system.
- State machines are a powerful approach to certain such services.
- A state machine
  - has a stored state, & receives inputs
  - Makes state transitions on each input & may output some results
    - \* transitions & output must be deterministic.
- A replicated state machine is a state machine that is replicated on multiple nodes
  - All replicas must get exactly same inputs.
  - Even if some of nodes fail, state & output can be obtained from other nodes

### Replicated State Machine



Leader declares log of nodes. Updates log of state machine at each replica happens only after log is replicated at a majority.

Record has been committed.

### Uses of Replicated state machines

- Inputs can specify operations with parameters but operations must be deterministic
- result of operation can be sent from any replica
  - gets executed only when log record is committed in replicated log.
  - usually sent from leader, which knows which part of log is committed.

Eg: Fault tolerant key-value store:-

- State: key-value storage state
- Operations: get() & put()
- operations executed when log is in committed state.

### Part - A

- 3) Advantages of storing multiple relations in a single file:
- Complex structures can be implemented through the DBMS, thus increasing performance.
  - Organise data & programs on storage media in a human friendly way.
  - Editing data is easier.
- Disadvantages:-
- Data backup: If the single file is corrupted, all data is lost, so no data backup.
  - Increases size & complexity of DBMS.
  - Data inconsistency

- 4) For the two disk mirrored case, we assume A disk & B disk. In order to lose data, A & B need to fail at same time. If A is already failed & within 100000 hrs B disk will fail, then data is lost. The other case is B is already failed & within 100000 hrs A will fail & then data will be lost.  
For first case, A disk is failed for 100 hrs every 100000 hrs. So in order to make B to fail, it will need  $100000^n / 100$  hrs. Because the other case, the time is reduced to  $100000^n / (2 * 100)$ .

- 5) Map database management systems are software programs designed to efficiently store & recall spatial information. They are widely used in localisation & navigation, especially in automotive applications. A map database enables the rapid indexing & lookup of a large amount of geographic data.
- 4) MySQL enables restrictions to be placed on reuse of previous passwords. MySQL has password reuse policy that allows it, to establish password-reuse policy globally, via password-history & password-reuse-interval system variables.
- 2) Database indexing. Hash tables may also be used as disk-based data structures & database indices although B-trees are more popular in these applications. In multi-node database systems, hash tables are commonly used to distribute rows amongst nodes, reducing network traffic for hash joins.